

Replicating Technical Specification Documents using Retrieval-Augmented Generation (RAG)

Name: Yoshita Banerjee

Guide: Prof. Senthilmurugan Subbiah

Dept: Chemical Engineering, BTP 1

Problem Statement

- Technical specification documents are dense, multimodal (texts, tables, figures), and domain specific
- Large language models (LLMs) often hallucinate or misinterpret acronyms/definitions
- Requirement: A pipeline that can reliably retrieve, structure and regenerate technical documentation

Literature Review

- **Vanilla RAG** (*Lewis et al. 2020*) : Retrieval + Generation, but weak for domain specific docs
- **RAG for technical documents** (*Soman and Roychowdhury, 2024*) : Found that sentence-level retrieval and splitting term-definition pairs work best
- **MA-RAG** (*Nguyen et al., 2025*): Multi-agent approach (Planner, Step-Definer, Extractor, QA Writer) for handling ambiguities
- **SAMAC-R3-MED** (*Li et al., 2025*): Multi-modal semantic trees (MMST) + knowledge graphs for engineering documents

Gaps to be solved

Existing systems do not focus on replicating structured specification documents

Need a prototype that can achieve the following tasks

**Parse sections,
acronyms, and
tables
(Dolphin/MMST
idea)**

**Retrieve relevant
chunks at
sentence-level**

**Use of GPT API
to generate
cleaner version
post extraction**

Proposed Solutions

Workflow for BTP-1

1. Parser- Extract headers, paragraphs, acronyms, tables
2. Retriever: Sentence-level retrieval (FAISS, SBERT embeddings)
3. Generator: GPT API returns regenerated section
4. Evaluation: Use simple metrics → cosine similarity, BLEU/ROUGE scores

Deliverable - A demo where upon giving an input query or section → model regenerates a spec-style output with inline citation

RETRIEVAL QUALITY

- Cosine similarity
- Recall@k

Evaluation Metrics

GENERATION QUALITY

- BLEU
- ROUGE

Scope of BTP-2

- Incorporate multi-agent MA-RAG workflow (Planner, Step-Definer, Extractor, QA Writer)
- Handle tables and figures using layout-aware parsing (Dolphin, MMST)
- Improve evaluation with expert feedback loops and better semantic metrics

Conclusion

- BTP-1: Build a demonstrable RAG prototype for specification document replication
- BTP-2: Extend to multi-agent + multimodal retrieval for robust work

References

RAG systems for Technical Documents

Multi-Agent RAG

Multi-Agent collaboration of R3 models for multi-modal documents

Dolphin (Document Image Parsing via Heterogenous Anchor Prompting